	Name	Released by	Model size, billion	Model
1	Aya	Cohere	13	aya-101
2	Viking	Silo AI & TurkuNLP & HPLT	7, 13 and 33	Viking-7B, Viking-13B, Viking-33B
	Falcon 2	Technology Innovation Institute	1.3, 7.5, 11, 40, 180	Falcon 11B, falcon-180B
3	Falcon2-5.5B-Norwegian		5.5	Falcon2-5.5B-Norwegian
	Mistral	Mistral AI team	7	
4	NorMistral-7b-warm	collaboration between the Language Technology Group at the University of Oslo, the High Performance Language Technologies (HPLT) project, the National Library of Norway, and the University of Turku	7	
5	NorMistral-7b-scratch	collaboration between the Language Technology Group at the University of Oslo, the High Performance Language Technologies (HPLT) project, the National Library of Norway, and the University of Turku	7	
6	Llama-2-13b-chat-norwegian	Ruter	13	

Aya from Cohere

About

Aya is one of the largest open science endeavors in machine learning to date – redefining the research landscape by collaborating with independent researchers from across the globe. The result is a fully **open-sourced** dataset and model.

The Aya Model is a massively multilingual language model capable of following instructions in 101 languages. Developed using a diverse mix of instructions from the Aya dataset and collection among others, it achieves state-of-the-art performance across numerous multilingual benchmarks.

Released

Feb 13, 2024

Viking from Silo AI

About

Silo AI is Europe's largest private AI lab. Together with University of Turku's research group TurkuNLP and HPLT, Silo AI releases the first checkpoints of Viking, an open LLM for all **Nordic** languages, English and programming languages.

Viking provides models that are not only linguistically performant and inclusive, but also sensitive to local values and cultures. Evaluations indicate best-in-class performance in all Nordic languages, without compromising performance in English. Viking is an enhanced model family of the previous model Poro, and is further evidence of the novel approach to train LLMs for low-resource languages

Released

Apr 4, 2024

The Falcon Series of Open Language Models

About

We introduce the Falcon series: 7B, 40B, and 180B parameters causal decoder- only models trained on a diverse high-quality corpora predominantly assembled from web data. The largest model, Falcon-180B, has been trained on over 3.5 trillion tokens of text–the largest openly documented pretraining run. Falcon-180B significantly outperforms models such as PaLM or Chinchilla, and improves upon concurrently developed models such as LLaMA 2 or Inflection-1. It nears the performance of PaLM-2-Large at a reduced pretraining and inference cost, making it, to our knowledge, one of the three best language models in the world along with GPT-4 and PaLM-2-Large. We report detailed evaluations, as well as a deep dive into the methods and custom tooling employed to pretrain Falcon. Notably, we report on our custom distributed training codebase, allowing us to efficiently pretrain these models on up to 4,096 A100s on cloud AWS infrastructure with limited interconnect. We release a 600B tokens extract of our web dataset, as well as the Falcon-7/40/180B models under a permissive license to foster open-science and accelerate the development of an open ecosystem of large language models

Released

Nov 29, 2023